

FALL 2018 CISC 311

OBJECT & STRUCTURE & ALGORITHM I

– Chapter 8: Priority Queues and Heaps




Outline

- The priority queue ADT
- The heap ADT
- Implementations
- Heap sort




The Priority Queue ADT

- A priority queue is a collection of prioritized elements that allows arbitrary element insertion, and allows the removal of the element that has first priority
 - Insertions are at arbitrary positions
 - Removal are at the front of the priority queue
- The priority queue ADT stores arbitrary objects
 - Example:



Priority Queue Methods

- `insert(e)`: inserts element e to the back of priority queue
 - Step 1: increment `END` and shift all the elements greater than e forward so it points to the appropriate space




Priority Queue Methods (cont'd.)

- `insert(e)`: inserts element *e* to the back of priority queue
 - Step 2: insert new object




Priority Queue Methods (cont'd.)

- `remove()` : removes and returns the first element (which has first priority) from the priority queue or `null` if the priority queue is empty
 - Step 1: remove the value at front



Priority Queue Methods (cont'd.)

- `remove()`: removes and returns the first element (which has first priority) from the priority queue or `null` if the priority queue is empty
 - Step 2: decrement `FRONT`



Priority Queue Methods (cont'd.)

Method	Description
<code>size ()</code>	Returns the number of elements in the priority queue
<code>isEmpty ()</code>	Returns a boolean indicating whether the priority queue is empty
<code>min ()</code>	Returns the first element of the priority queue, without removing it or <code>null</code> if the priority queue is empty
<code>insert (e)</code>	Adds element <code>e</code> to priority queue
<code>remove ()</code>	Removes and returns the first element from the priority queue or <code>null</code> if the priority queue is empty




Priority Queue Implementation

- Array-based implementation (sorted list)
- Doubly linked list-based implementation (unsorted & sorted lists)
- Heap-based implementation




Sequence-based Priority Queue

- Implementation with a sorted list
- Performance:
 - `insert` takes $O(n)$ time since we have to find the place where to insert the item
 - `remove` and `min` take $O(1)$ time, since the smallest key is at the beginning
- Implementation with an unsorted list
- Performance:
 - `insert` takes $O(1)$ time since we can insert the item at the beginning or end of the sequence
 - `remove` and `min` take $O(n)$ time since we have to traverse the entire sequence to find the smallest key



Array-based Implementation

- Use an array of size N
- Two variables keep track of the front and size
 - m : index of the front element
 - sz : number of stored elements



Array-based Implementation (cont'd.)

Algorithm `size()`

return `sz`

Algorithm `isEmpty()`


return `sz == 0`

Algorithm `min()`

if `isEmpty()` then

return `null`

return `Q[m]`




`ArrayPriorityQueue.java`



Array-based Implementation (cont'd.)

- Note that operation `remove` returns `null` if the priority queue is empty

```
Algorithm remove()  
  if isEmpty() then  
    return null  
  removed  $\leftarrow Q[m]$   
  m  $\leftarrow m + 1$   
  sz  $\leftarrow sz - 1$   
  return removed
```




ArrayPriorityQueue.java



Array-based Implementation (cont'd.)

- Operation `insert` throws an exception if the array is full

```
Algorithm insert(o)  
  if sz = Q.length then  
    throw IllegalStateException  
  r ← m + sz  
  for i ← r - 1 to 0 do  
    if o < S[i] then  
      S[i+1] ← S[i]  
    else  
      break  
  Q[i+1] ← o  
  sz ← sz + 1
```



ArrayPriorityQueue.java



Total Order Relations

- Keys in a priority queue can be arbitrary objects on which an order is defined
- Two distinct entries in a priority queue can have the same key
- Mathematical concept of total order relation \leq
 - Comparability property: either $x \leq y$ or $y \leq x$
 - Antisymmetric property: $x \leq y$ and $y \leq x \Rightarrow x = y$
 - Transitive property: $x \leq y$ and $y \leq z \Rightarrow x \leq z$
 - Reflexive property: $k \leq k$



Comparator Interface

- A comparator encapsulates the action of comparing two objects according to a given total order relation
- A generic priority queue uses an auxiliary comparator
- The comparator is external to the keys being compared
- When the priority queue needs to compare two keys, it uses its comparator
- Implement `java.util.Comparator`
- Primary method of Comparator
 - `compare(a, b)`: returns an integer i such that
 - $i < 0$ if $a < b$
 - $i = 0$ if $a = b$
 - $i > 0$ if $a > b$
 - An error occurs if a and b cannot be compared



Comparable Interface

- A comparable encapsulates the action of comparing two objects according to the natural ordering
- Implement `java.lang.Comparable`
- It only contains one method
 - `a.compareTo(b)`: returns an integer i such that
 - $i < 0$ if $a < b$
 - $i = 0$ if $a = b$
 - $i > 0$ if $a > b$
 - An error occurs if a and b cannot be compared
- E.g., the natural ordering of strings is lexicographic, which is a case-sensitive extension of the alphabetic ordering of Unicode

`StringLengthComparator.java`



Performance Comparison

Method	Unsorted linked list	Sorted linked list
<code>size()</code>	$O(1)$	$O(1)$
<code>isEmpty()</code>	$O(1)$	$O(1)$
<code>min()</code>	$O(n)$	$O(1)$
<code>insert(e)</code>	$O(1)$	$O(n)$
<code>remove()</code>	$O(n)$	$O(1)$

`UnsortedPriorityQueue.java`

`SortedPriorityQueue.java`



Priority Queue ADT (cont'd.)

- A priority queue stores a collection of entries
- Each entry is a pair (*key*, *value*)
- Main methods of the priority queue ADT
 - `insert(k, v)`: inserts an entry with key *k* and value *v*
 - `remove()`: removes and returns the entry with smallest key, or `null` if the the priority queue is empty
- Additional methods
 - `min()`: returns, but does not remove, an entry with smallest key, or `null` if the priority queue is empty
 - `size()`
 - `isEmpty()`

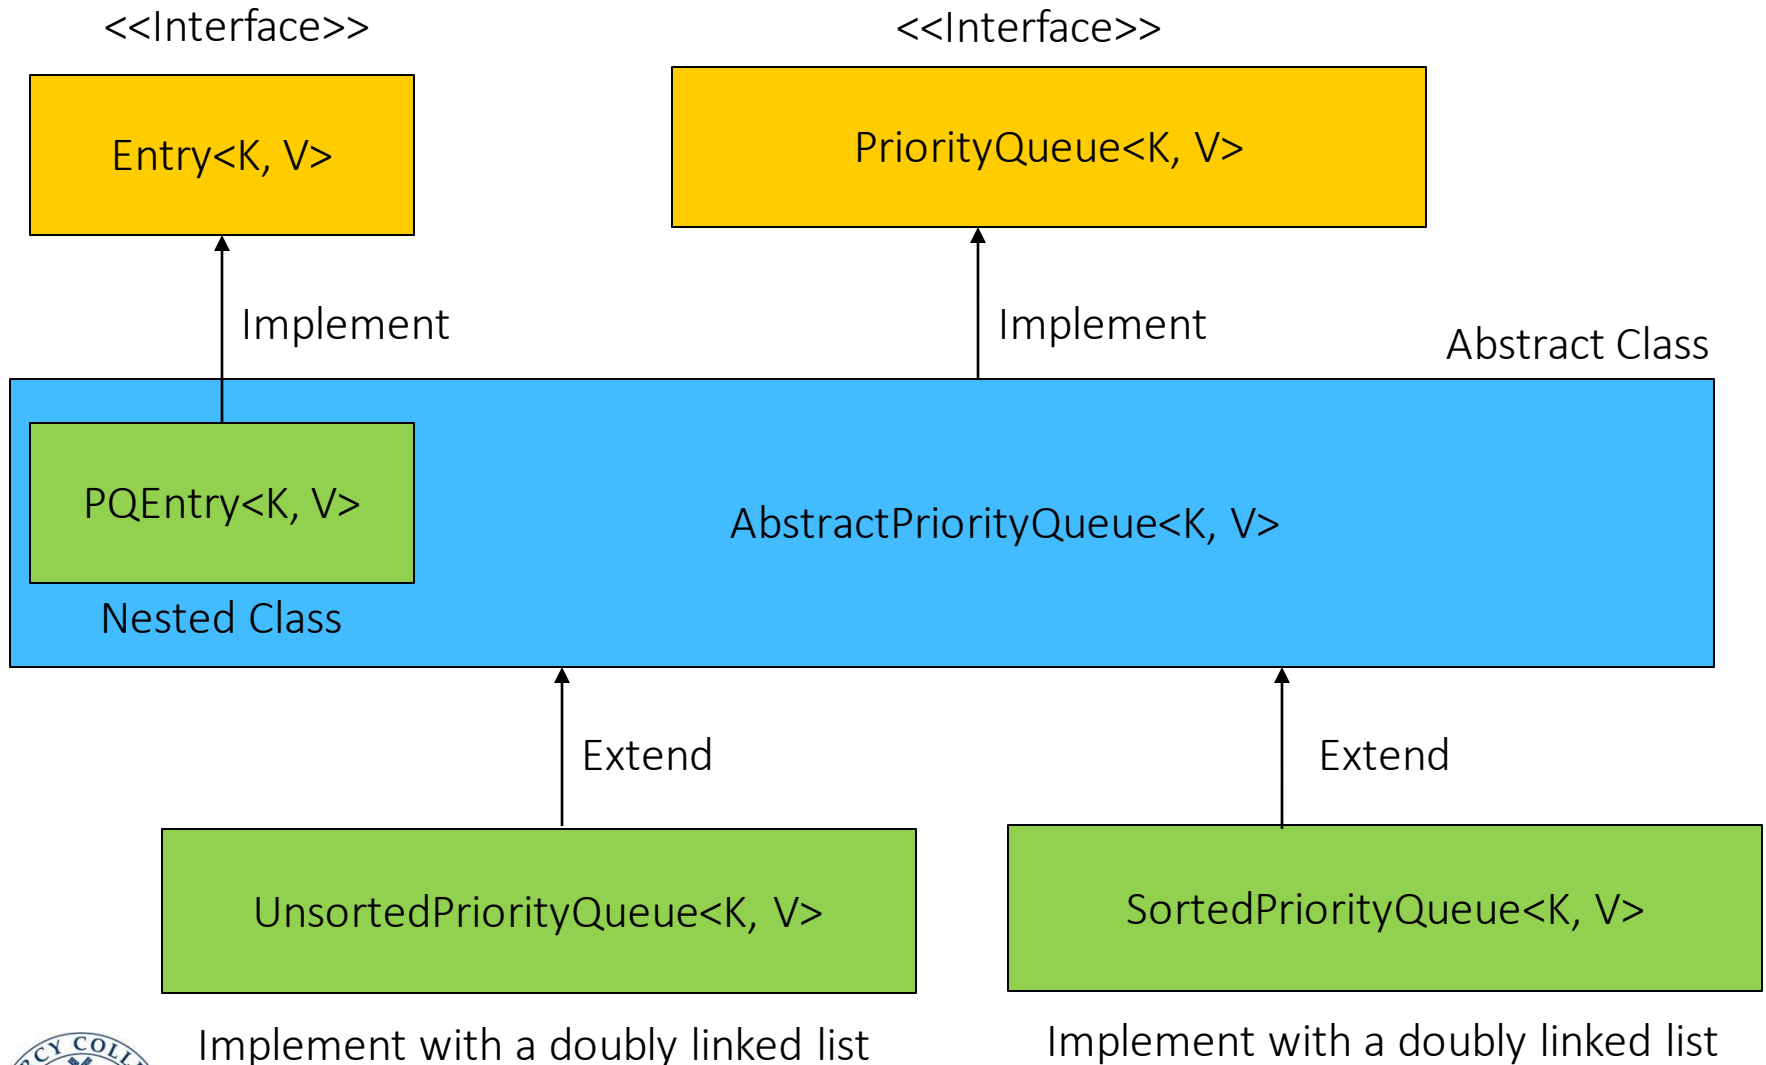


Priority Queue Methods (cont'd.)

Method	Description
<code>size()</code>	Returns the number of elements in the priority queue
<code>isEmpty()</code>	Returns a boolean indicating whether the priority queue is empty
<code>min()</code>	Returns a priority queue entry (k, v) having the minimal key, without removing it or <code>null</code> if the priority queue is empty
<code>insert(k, v)</code>	Inserts an entry with key k and value v in the priority queue
<code>remove()</code>	Removes and returns an entry (k, v) having minimal key from the priority queue or <code>null</code> if the priority queue is empty



Implementation



Implement with a doubly linked list

Implement with a doubly linked list

Entry ADT

- An entry in a priority queue is simply a key-value pair
- Priority queues store entries to allow for efficient insertion and removal based on keys
- Methods:
 - `getKey()`: returns the key for this entry
 - `getValue()`: returns the value associated with this entry


`Entry.java`

`PQEntry.java`




The Heap ADT

- A heap is a binary tree storing keys at its nodes and satisfying the following properties:
 - Heap-Order: for every internal node v other than the root, $key(v) \geq key(parent(v))$
 - Complete Binary Tree: let h be the height of the heap
 - For $i = 0, \dots, h - 1$, there are 2^i nodes of depth i
 - At depth $h - 1$, the internal nodes are to the left of the external nodes
 - The last node of a heap is the rightmost node of maximum depth




Height of a Heap

- Theorem: A heap storing n keys has height $O(\log n)$
- Proof: (we apply the complete binary tree property)
 - Let h be the height of a heap storing n keys
 - Since there are 2^i keys at depth $i = 0, \dots, h-1$ and at least one key at depth h , we have $n \geq 1 + 2 + 4 + \dots + 2^{h-1} + 1$
 - Thus, $n \geq 2^h$, i.e., $h \leq \log n$





Heaps and Priority Queues

- We can use a heap to implement a priority queue
- We store a $(key, element)$ item at each internal node
- We keep track of the position of the last node




Insertion into a Heap

- Method `insert` of the priority queue ADT corresponds to the insertion of a key k to the heap
- The insertion algorithm consists of three steps
 - Find the insertion node z (the new last node)
 - Store k at z
 - Restore the heap-order property





Upheap

- After the insertion of a new key k , the heap-order property may be violated
- Algorithm upheap restores the heap-order property by swapping k along an upward path from the insertion node
- Upheap terminates when the key k reaches the root or a node whose parent has a key smaller than or equal to k
- Since a heap has height $O(\log n)$, upheap runs in $O(\log n)$ time





Removal from a Heap

- Method `remove` of the priority queue ADT corresponds to the removal of the root key from the heap
- The removal algorithm consists of three steps
 - Replace the root key with the key of the last node w
 - Remove w
 - Restore the heap-order property




Downheap

- After replacing the root key with the key k of the last node, the heap-order property may be violated
- Algorithm downheap restores the heap-order property by swapping key k along a downward path from the root
- Upheap terminates when key k reaches a leaf or a node whose children have keys greater than or equal to k
- Since a heap has height $O(\log n)$, downheap runs in $O(\log n)$ time

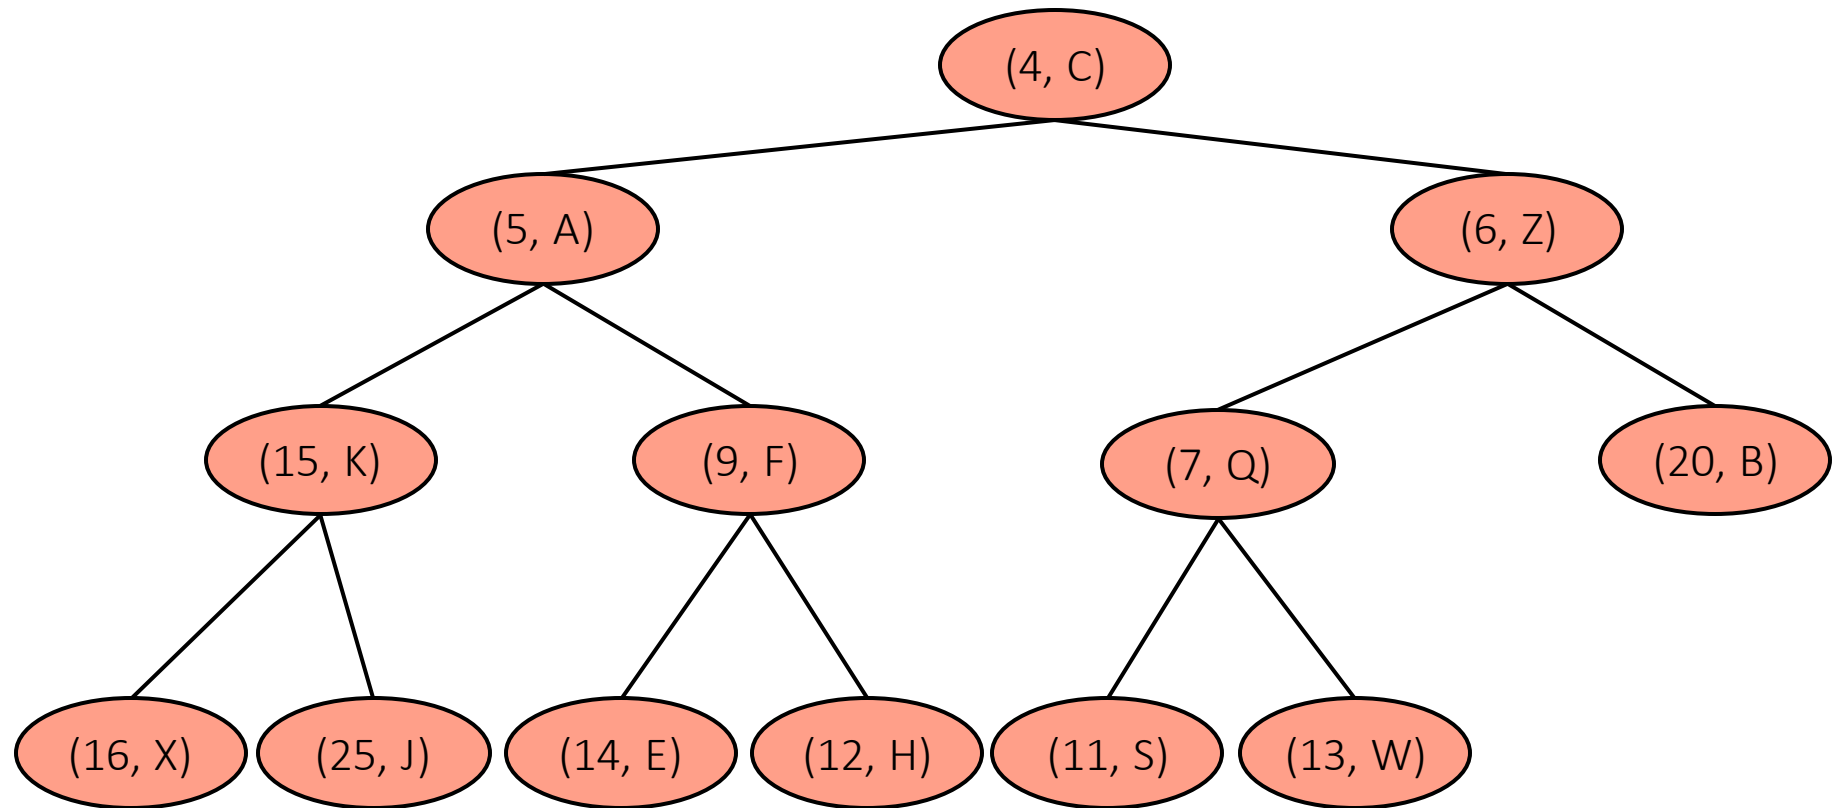


Updating the Last Node

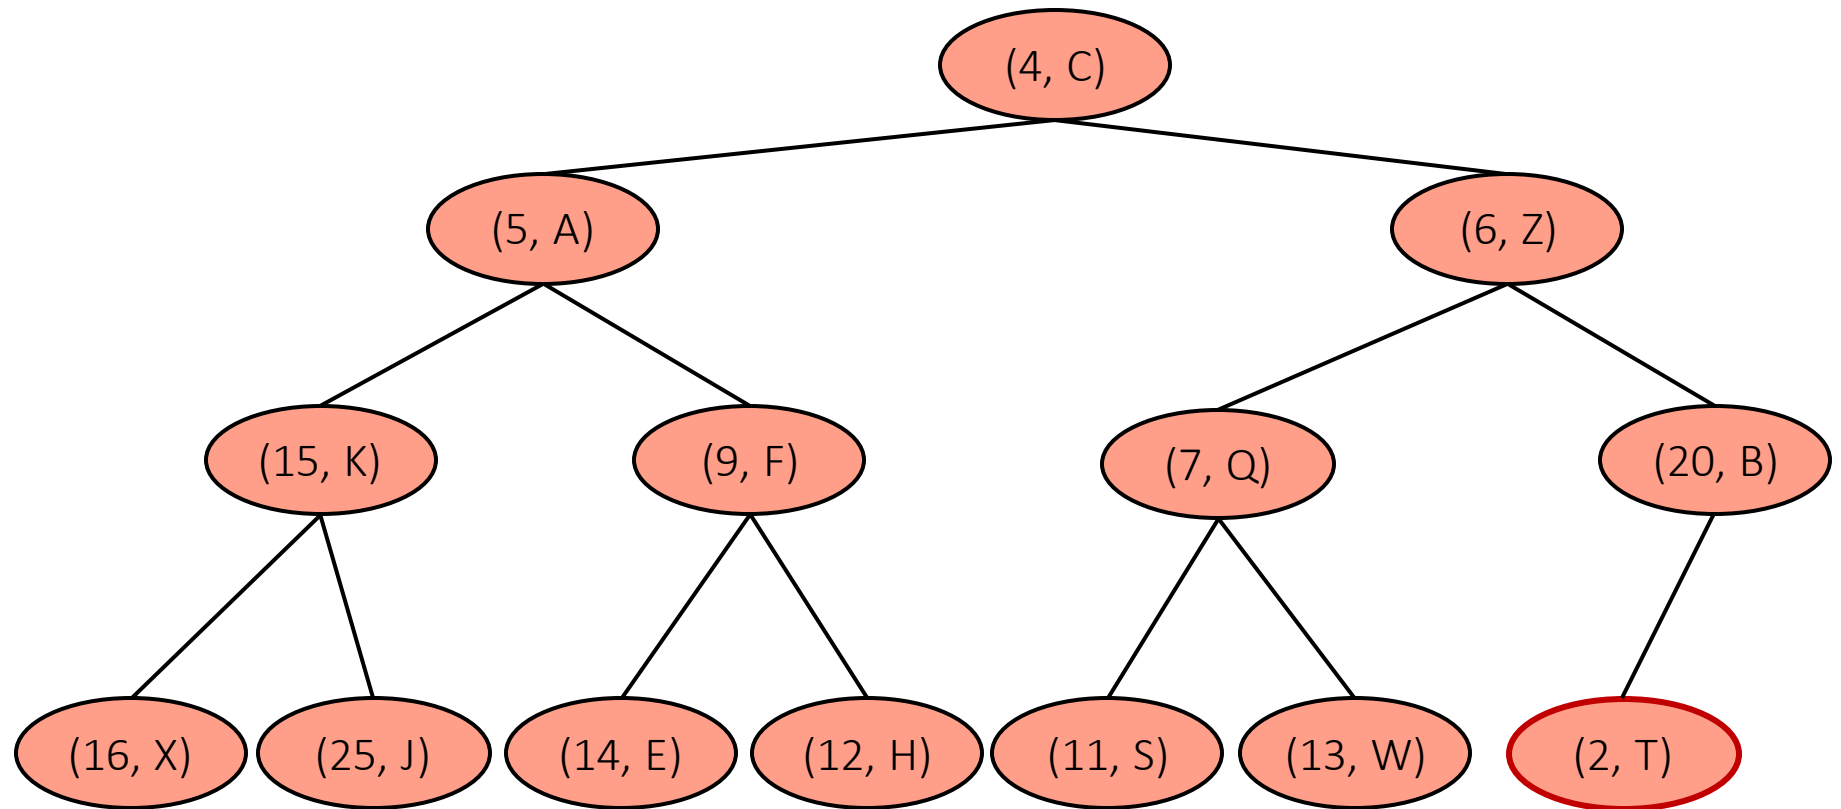
- The insertion node can be found by traversing a path of $O(\log n)$ nodes
 - Go up until a left child or the root is reached
 - If a left child is reached, go to the right child
 - Go down left until a leaf is reached
- Similar algorithm for updating the last node after a removal




Heap-based Implementation




Upheap



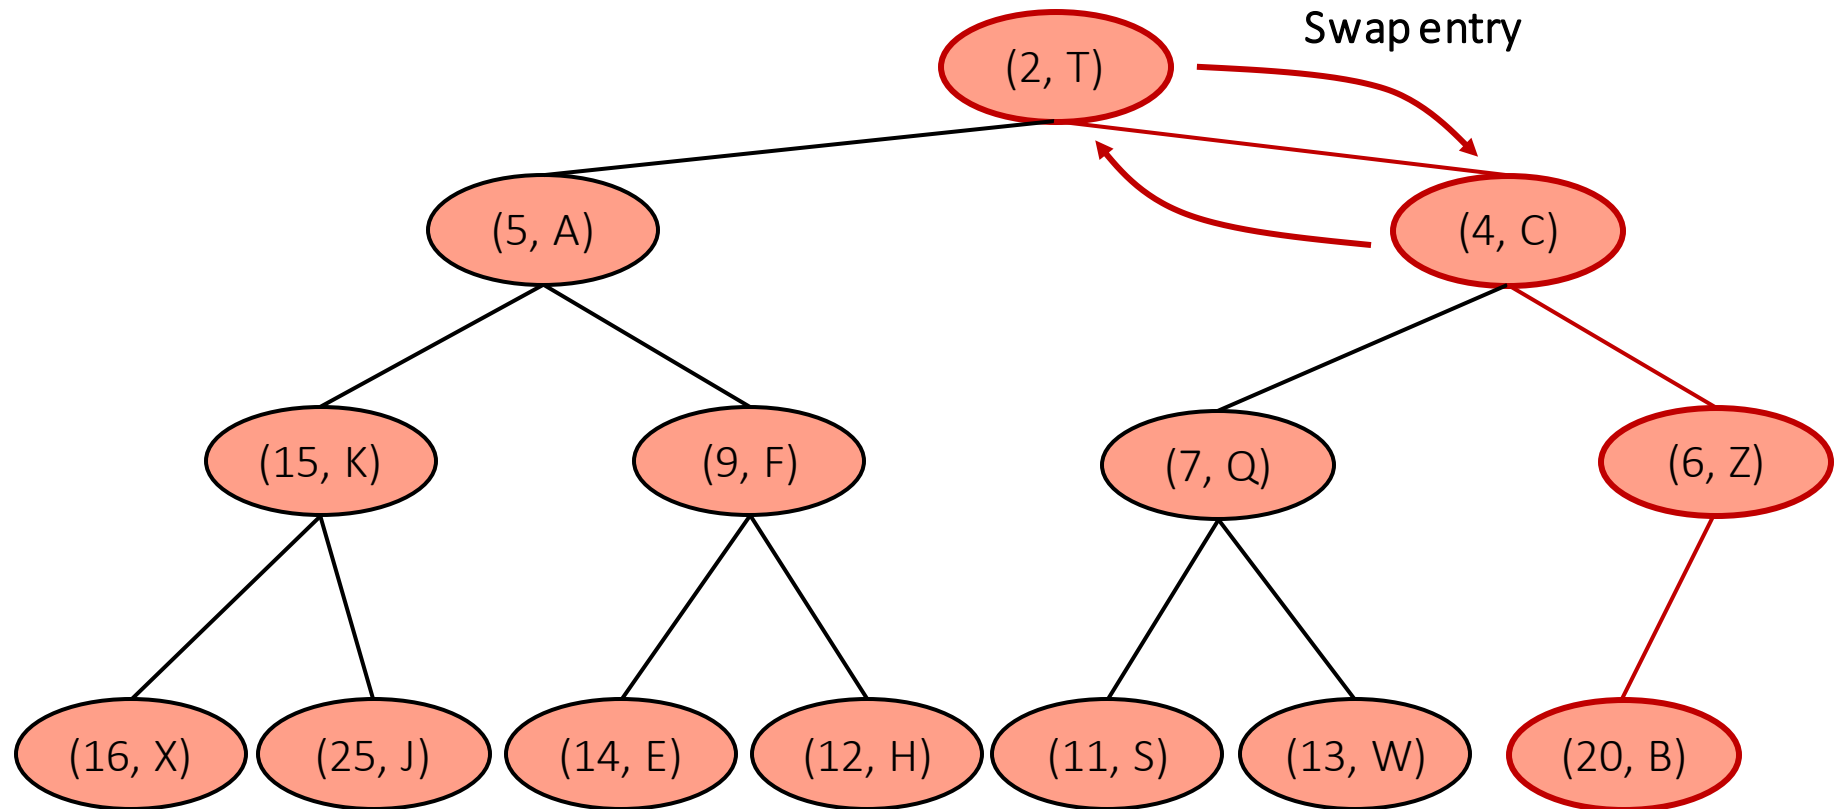
Upheap (cont'd.)



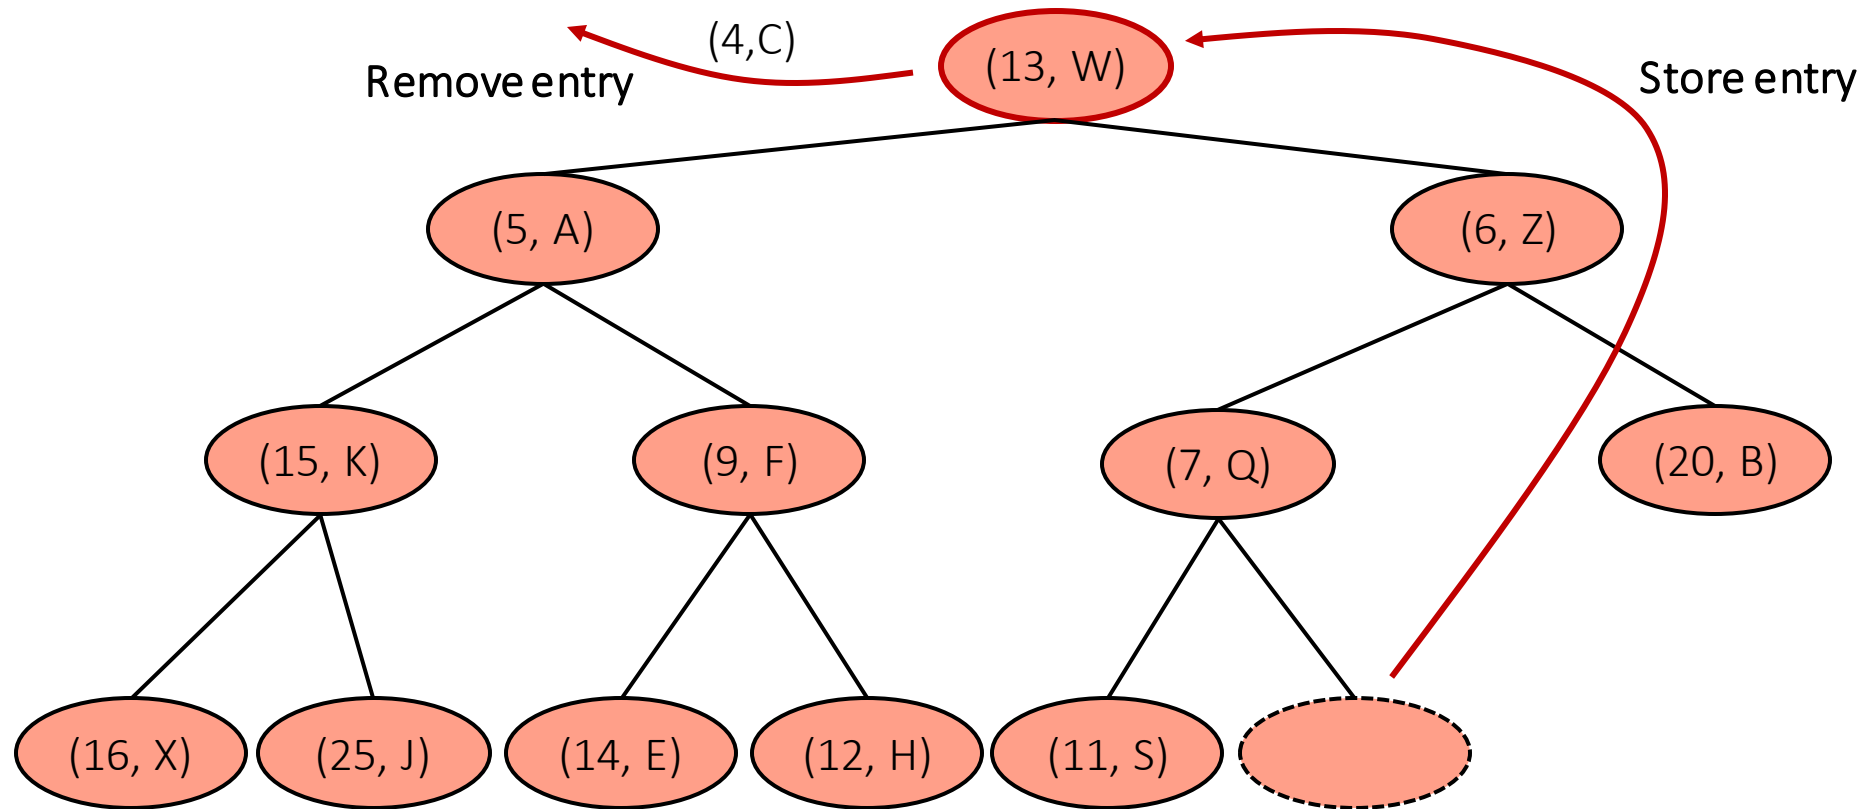
Upheap (cont'd.)




Upheap (cont'd.)




Downheap



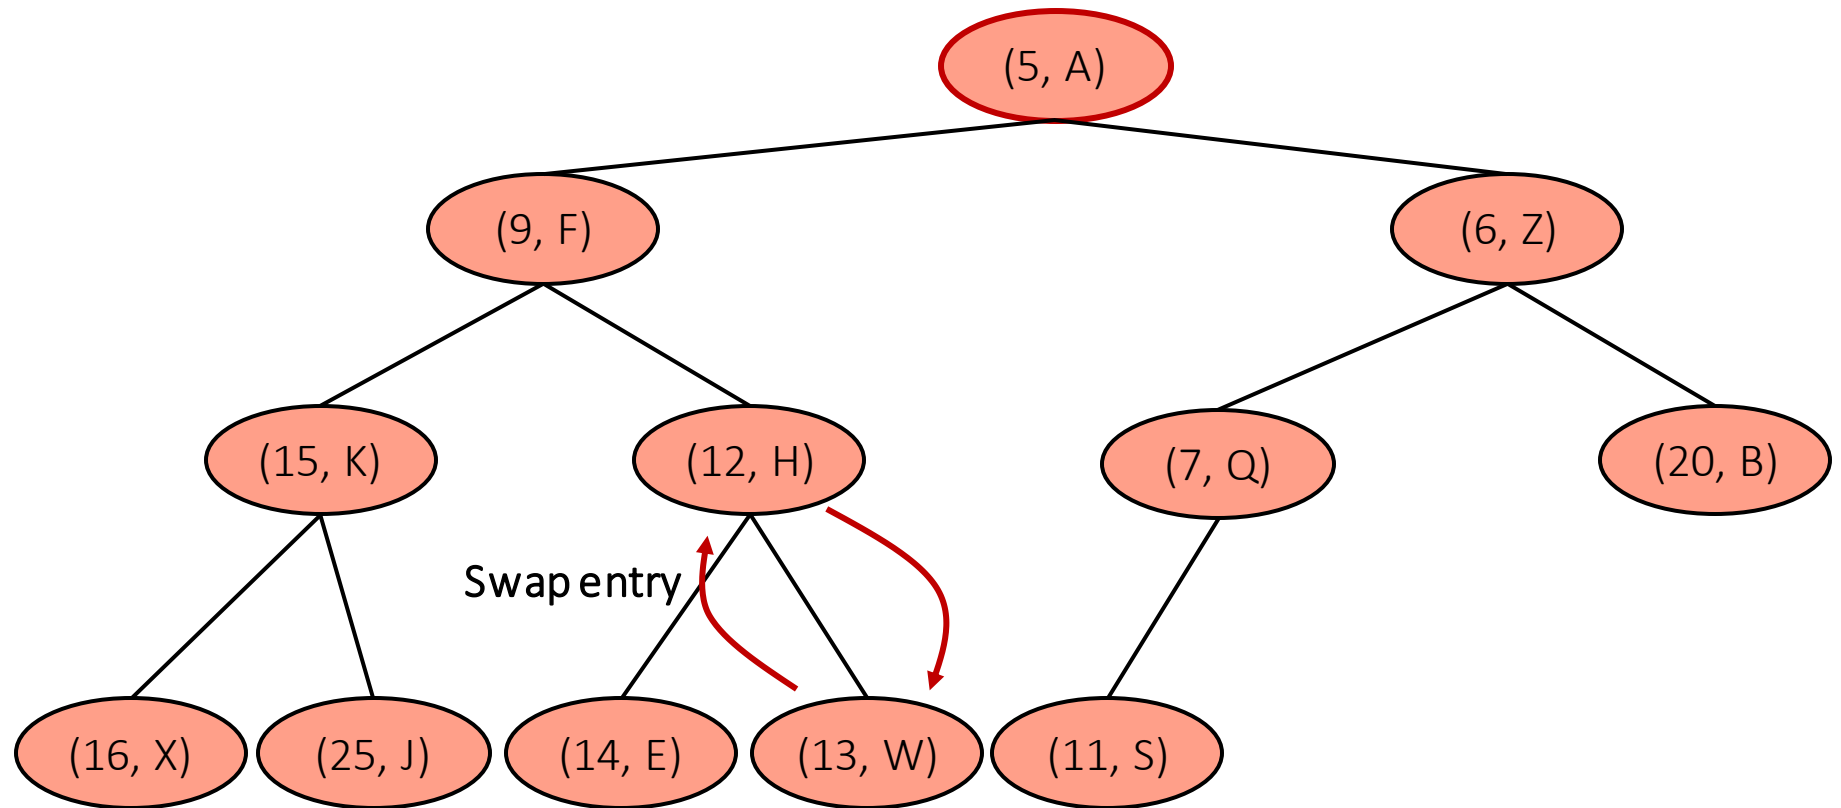
Downheap (cont'd.)



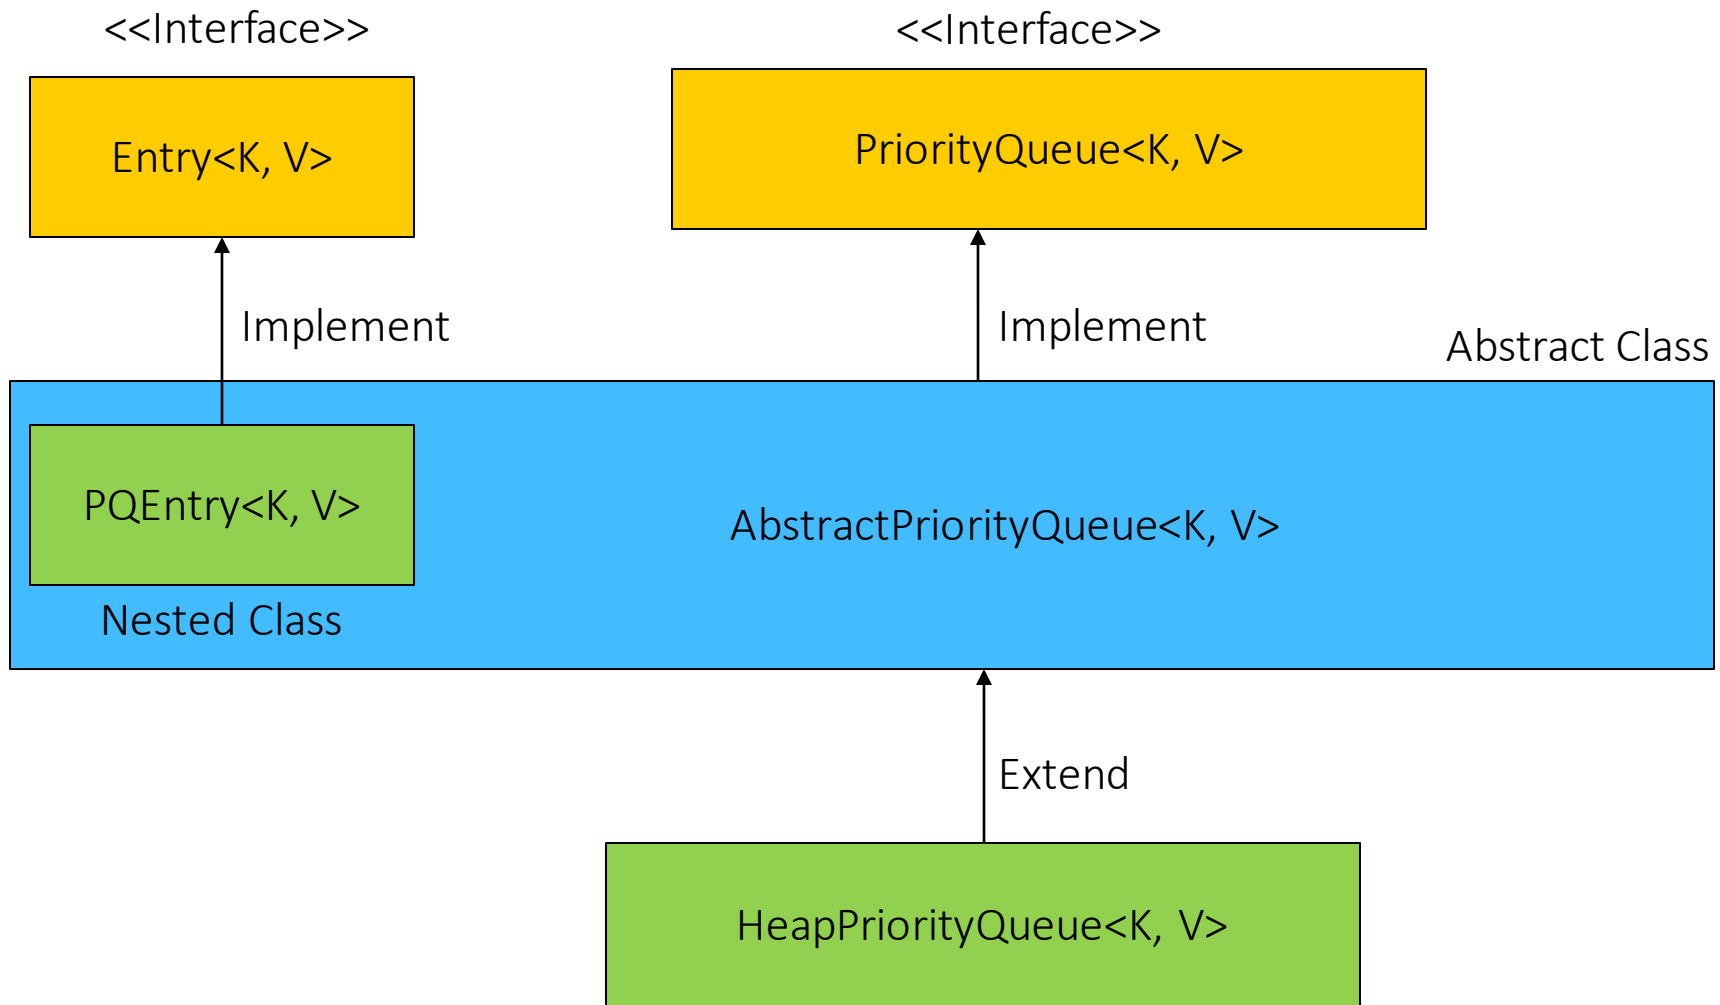
Downheap (cont'd.)



Downheap (cont'd.)



Implementation



Recall Array-Based Binary Trees

- Based on a way of numbering the node of T
 - If p is the root of T , then $f(p) = 0$
 - If p is the left child of node q , then $f(p) = 2f(q) + 1$
 - If p is the right child of node q , then $f(p) = 2f(q) + 2$
 - If p is the parent of node q , then $f(p) = (f(q) - 1)/2$
 - $f()$: Level numbering of the nodes in a binary tree T

`ArrayBinaryTree.java`





Recall Array-Based Binary Trees (cont'd.)

- Nodes are stored in an array A




- Node v is stored at $A[\text{rank}(v)]$
 - $\text{rank}(\text{root}) = 0$
 - If node is the left child of $\text{parent}(\text{node})$,
 $\text{rank}(\text{node}) = 2 \cdot \text{rank}(\text{parent}(\text{node})) + 1$
 - If node is the right child of $\text{parent}(\text{node})$,
 $\text{rank}(\text{node}) = 2 \cdot \text{rank}(\text{parent}(\text{node})) + 2$





Array-based Heap Implementation

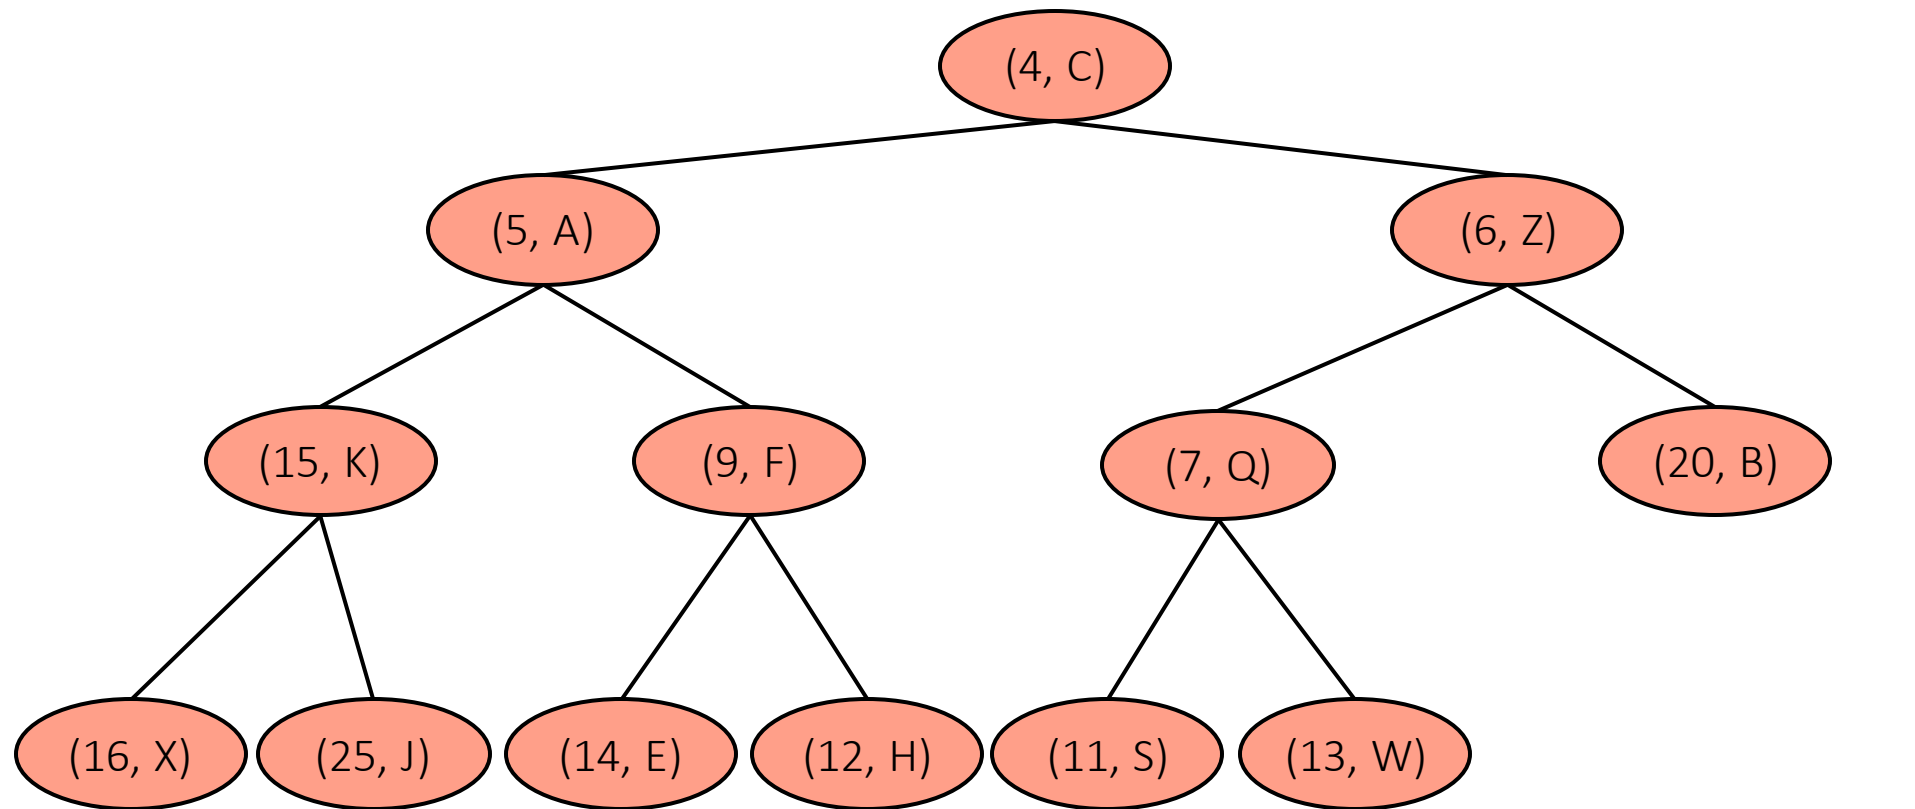
- We can represent a heap with n keys by means of an array of length n
- For the node at rank i
 - The left child is at rank $2i + 1$
 - The right child is at rank $2i + 2$
- Links between nodes are not explicitly stored
- Operation add corresponds to inserting at rank $n + 1$
- Operation remove corresponds to removing at rank n
- Yields in-place heap-sort



2	5	6	9	7
0	1	2	3	4

`HeapPriorityQueue.java`





(4, C)	(5, A)	(6, Z)	(15, K)	(9, F)	(7, Q)	(20, B)	(16, X)	(25, J)	(14, E)	(12, H)	(11, S)	(13, W)
--------	--------	--------	---------	--------	--------	---------	---------	---------	---------	---------	---------	---------

0 1 2 3 4 5 6 7 8 9 10 11 12



Performance Comparison of Priority Queue

Method	Sorted array	Unsorted linked list	Sorted linked list	Heap
<code>size()</code>	$O(1)$	$O(1)$	$O(1)$	$O(1)$
<code>isEmpty()</code>	$O(1)$	$O(1)$	$O(1)$	$O(1)$
<code>min()</code>	$O(1)$	$O(n)$	$O(1)$	$O(1)$
<code>insert(e)</code>	$O(n)$	$O(1)$	$O(n)$	$O(\log n)$
<code>remove()</code>	$O(1)$	$O(n)$	$O(1)$	$O(\log n)$



Priority Queue Sorting

- We can use a priority queue to sort a list of comparable elements
 - Insert the elements one by one with a series of insert operations
 - Remove the elements in sorted order with a series of remove operations

Algorithm PQ-Sort(S , C)

Input list S , comparator C for the elements of S

Output list S sorted in increasing order according to C

$P \leftarrow$ priority queue with comparator C

while $\neg S.\text{isEmpty}()$ **do**

$e \leftarrow S.\text{remove}(S.\text{first}())$

$P.\text{insert}(e, \emptyset)$

while $\neg P.\text{isEmpty}()$ **do**

$e \leftarrow P.\text{remove}().\text{getKey}()$

$S.\text{addLast}(e)$



Priority Queue Sorting

- The running time depends on the priority queue implementation:
 - Unsorted sequence gives selection-sort: $O(n^2)$ time
 - Sorted sequence gives insertion-sort: $O(n^2)$ time
- Can we do better?

Algorithm PQ-Sort(S , C)

Input list S , comparator C for the elements of S

Output list S sorted in increasing order according to C

$P \leftarrow$ priority queue with comparator C

while $\neg S.\text{isEmpty}()$ **do**

$e \leftarrow S.\text{remove}(S.\text{first}())$

$P.\text{insert}(e, \emptyset)$

while $\neg P.\text{isEmpty}()$ **do**

$e \leftarrow P.\text{remove}().\text{getKey}()$

$S.\text{addLast}(e)$



Heap-Sort

- Consider a priority queue with n items implemented by means of a heap
 - The space used is $O(n)$
 - Methods insert and remove take $O(\log n)$ time
 - Methods `size()`, `isEmpty()`, and `min()` take time $O(1)$ time
- Using a heap-based priority queue, we can sort a sequence of n elements in $O(n \log n)$ time
- The resulting algorithm is called heap-sort
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

HeapSort.java



Recall Selection Sort

- Selection-sort is the variation of PQ-sort where the priority queue is implemented with an unsorted sequence
- Running time of selection-sort:
 - Inserting the elements into the priority queue with n insert operations takes $O(n)$ time
 - Removing the elements in sorted order from the priority queue with n remove operations takes time proportional to
$$1 + 2 + \dots + n$$
- Selection-sort runs in $O(n^2)$ time



Selection Sort Example

	Sequence S	Priority Queue P
Input:	(7,4,8,2,5,3,9)	()
Phase 1		
(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(7,4)
..	
(g)	()	(7,4,8,2,5,3,9)
Phase 2		
(a)	(2)	(7,4,8,5,3,9)
(b)	(2,3)	(7,4,8,5,9)
(c)	(2,3,4)	(7,8,5,9)
(d)	(2,3,4,5)	(7,8,9)
(e)	(2,3,4,5,7)	(8,9)
(f)	(2,3,4,5,7,8)	(9)
(g)	(2,3,4,5,7,8,9)	()



Recall Insertion Sort

- Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted sequence
- Running time of Insertion-sort:
 - Inserting the elements into the priority queue with n insert operations takes time proportional to
$$1 + 2 + \dots + n$$
 - Removing the elements in sorted order from the priority queue with a series of n remove operations takes $O(n)$ time
- Insertion-sort runs in $O(n^2)$ time




Insertion-Sort Example

	Sequence S	Priority queue P
Input:	(7,4,8,2,5,3,9)	()
Phase 1		
(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(4,7)
(c)	(2,5,3,9)	(4,7,8)
(d)	(5,3,9)	(2,4,7,8)
(e)	(3,9)	(2,4,5,7,8)
(f)	(9)	(2,3,4,5,7,8)
(g)	()	(2,3,4,5,7,8,9)
Phase 2		
(a)	(2)	(3,4,5,7,8,9)
(b)	(2,3)	(4,5,7,8,9)
..
(g)	(2,3,4,5,7,8,9)	()



In-place Insertion-Sort

- Instead of using an external data structure, we can implement selection-sort and insertion-sort in-place
- A portion of the input sequence itself serves as the priority queue
- For in-place insertion-sort
 - We keep sorted the initial portion of the sequence
 - We can use swaps instead of modifying the sequence



Algorithm	Time	Notes
Selection sort	$O(n^2)$	<ul style="list-style-type: none"> in-place slow (good for small inputs) stable
Insertion sort	$O(n^2)$	<ul style="list-style-type: none"> in-place slow (good for small inputs) stable
Heap sort	$O(n \log n)$	<ul style="list-style-type: none"> in-place not stable
Quick sort	$O(n \log n)$ expected	<ul style="list-style-type: none"> in-place, randomized fastest (good for large inputs) not stable
Merge sort	$O(n \log n)$	<ul style="list-style-type: none"> fast sequential data access for huge data sets ($> 1M$) stable



Applications of Priority Queue

- Standby flyers
- Auctions
- Stock market



Summary

- Priority Queue ADT
 - Implementation:
 - Sorted array
 - Sorted doubly linked list
 - Unsorted doubly linked list
 - Heap
- Heap ADT: complete binary tree
 - Upheap, downheap
 - Implementation:
 - Array-list-based binary tree
- PQ sort
 - Heap sort

